

Lock-Free Ring Buffer for Market Data

Low Latency Data Pipeline Team

April 7, 2025

Outline

Introduction

Ring Buffer Design

Testing


Applications & Conclusions

Why Lock-Free Data Structures?

- ▶ **Challenge:** Modern market data systems process millions of updates per second
- ▶ **Problem:** Traditional mutex-based synchronization creates bottlenecks
- ▶ **Solution:** Lock-free data structures provide:
 - ▶ No mutex acquisition/release overhead
 - ▶ No thread blocking
 - ▶ Lower and more predictable latency
 - ▶ Better scalability under contention
- ▶ **Application:** Ideal for producer-consumer patterns in market data processing

Ring Buffer Fundamentals

- ▶ Fixed-size circular buffer
- ▶ Two atomic indices:
 - ▶ `write_idx_` - Next position to write
 - ▶ `read_idx_` - Next position to read
- ▶ Empty when `read_idx_ == write_idx_`
- ▶ Full when $(\text{write_idx_} + 1) \% \text{Size} == \text{read_idx_}$
- ▶ One slot always reserved for empty detection



Implementation Highlights

```
template <typename T, size_t Size>
class LockFreeRingBuffer {
public:
    bool TryPush(const T& item) {
        const size_t current_write = write_idx_.load(std::
memory_order_relaxed);
        const size_t next_write = (current_write + 1) % Size
;

        if (next_write == read_idx_.load(std::
memory_order_acquire))
            return false; // Buffer full

        buffer_[current_write] = item;
        write_idx_.store(next_write, std::
memory_order_release);
        return true;
    }

    bool TryPop(T* output) { /* Similar implementation */
    }
```

Key Technical Features

- ▶ **Lock-free algorithm:** Using atomic variables
- ▶ **Memory ordering:** Careful use of memory ordering constraints
 - ▶ `memory_order_relaxed` for initial loads
 - ▶ `memory_order_acquire` for checking conditions
 - ▶ `memory_order_release` for committing updates
- ▶ **Cache line alignment:** Preventing false sharing between indices
 - ▶ Producer and consumer threads operate on different cache lines
 - ▶ `alignas(64)` ensures indices are on separate cache lines
- ▶ **Non-blocking behavior:** `TryPush/TryPop` never block, return boolean success

Basic Functionality Test

- ▶ Tests fundamental properties with realistic market data structures:

```
struct MarketTick {  
    int64_t timestamp_ns;  
    std::string symbol;  
    double price;  
    double quantity;  
    char side; // 'B' for buy, 'S' for sell  
  
    // For testing equality in our assertions  
    bool operator==(const MarketTick(&@&@*) other)  
    const;  
};
```

- ▶ Verifies:
 - ▶ Empty buffer detection
 - ▶ Push until full
 - ▶ Pop in FIFO order
 - ▶ Alternating push/pop sequences


Market Data Pipeline Test

- ▶ Simulates realistic market data processing:
 - ▶ Producer thread: generates synthetic market tick data
 - ▶ Consumer thread: processes ticks and measures latency
- ▶ Test parameters:
 - ▶ 5,000 market ticks
 - ▶ 1024-slot buffer
 - ▶ Alternating BTC/ETH symbols
 - ▶ Realistic price movements
- ▶ Measures end-to-end latency from tick creation to processing

Performance Results

- ▶ **Test machine:** Modern x86_64 system
- ▶ **Compiler:** GCC 13.3.0
- ▶ **C++ standard:** C++23
- ▶ **Build type:** Release (-O3)
- ▶ **Minimum latency:** 215 ns
- ▶ **Median latency:** 26,305 ns
- ▶ **99th percentile:** 3,284,727 ns
- ▶ **Maximum latency:** 4,554,931 ns

Percentile



Applications in Market Data Systems

- ▶ **Feed handlers:** Buffering incoming market data packets
- ▶ **Parsing pipeline:** Moving raw data to normalization stage
- ▶ **Order book updates:** Efficiently queuing price updates
- ▶ **Strategy components:** Communicating signals between modules
- ▶ **Risk checks:** Queuing pre-trade validations
- ▶ **Logging:** Non-blocking capture of events for later analysis

Conclusions & Next Steps

► **Achievements:**

- Created a fully lock-free, high-performance ring buffer
- Demonstrated functionality with real market data structures
- Measured sub-microsecond minimum latency
- Successfully processed thousands of events with predictable performance

► **Potential improvements:**

- Multi-producer/multi-consumer variants
- Batched operations for higher throughput
- Memory reclamation for dynamically allocated elements
- Integration with hardware acceleration (FPGA, GPU)

Thank You

Questions?

Our implementation is available in the project repository:
`src/core/ring_buffer.h`